ZAK WHATMOUGH | EQUITY RESEARCH MARCH 2026

# NVIDIA CORPORATION

NVDA US EQUITY

NASDAQ | AI & GPU INFRASTRUCTURE

| RECOMMENDATION | 12-MONTH PRICE TARGET | CURRENT PRICE | IMPLIED UPSIDE |
|---|---|---|---|
| BUY | $213 | $178.03 | +19.6% |
| Based on 25x FY2027E EPS | FY2027E Non-GAAP EPS: $8.52 | 52-Wk Range: $86.62–$212.19 | vs. $272 street consensus |

| MKT CAP | EV | FY2026A REVENUE | FY2026A FCF MARGIN | NTM P/E | SECTOR |
|---|---|---|---|---|---|
| $4.46T | $4.42T | $215.9B | ~50% | 21.6x | Semiconductors |

## EXECUTIVE SUMMARY

We initiate coverage of NVIDIA Corporation (NVDA US) with a BUY recommendation and 12-month price target of $213, implying 19.6% upside from the current price of $178.03. Our target is derived from 25x FY2027E non-GAAP EPS of $8.52 — a deliberate discount to the Street consensus of $272 — reflecting our view that while NVIDIA's structural position is dominant, the stock is not a screaming buy at current levels. The 22% correction from its 52-week high has created a more attractive entry point, and the Rubin architecture ramp in FY2027 provides a credible near-term catalyst.

NVIDIA has transformed from a gaming GPU company into the defining infrastructure vendor of the AI era. FY2026 revenue reached $215.9 billion (+65% YoY), driven almost entirely by Data Center GPU demand. Gross margins have expanded to ~74% and non-GAAP operating margins exceed 65%. Free cash flow of $108 billion in FY2026 positions the company to return capital aggressively whilst maintaining investment in the next product cycle.

The investment case rests on three pillars: (1) the CUDA software moat, which creates genuine switching costs and has compounded for a decade; (2) the Blackwell-to-Rubin product cycle, with 3.6 million Blackwell units on backlog and sold out through mid-2026; and (3) structural hyperscaler capex commitments — Google ($180B), Amazon ($200B), and Microsoft (+80%) — that de-risk near-term demand visibility.

We are not at consensus. Our DCF at a 13% WACC produces an intrinsic value of $89, meaningfully below market. We acknowledge this and are explicit about it: our BUY is driven by forward multiples and product-cycle momentum, not DCF intrinsic value. The stock requires you to accept that NVIDIA commands a structural premium above DCF fair value — a view we find defensible for a near-monopoly position, but one that carries real downside risk should demand assumptions be revised.

Key risks include: (1) China export controls threatening ~13% of revenue; (2) DeepSeek-class efficiency improvements reducing per-query GPU demand; (3) hyperscaler custom ASICs (Google TPU v7, AWS Trainium) eroding market share at the margin; and (4) valuation compression if growth decelerates faster than expected.

## 1. BUSINESS OVERVIEW

### Company Description

NVIDIA Corporation designs and sells graphics processing units (GPUs), system-on-chip units, and software platforms for computing-intensive workloads. Founded in 1993 and headquartered in Santa Clara, California, NVIDIA employs approximately 36,000 people globally and operates on a fabless manufacturing model, outsourcing chip production primarily to Taiwan Semiconductor Manufacturing Company (TSMC).

The company reports across five segments: Data Center, Gaming, Professional Visualization, Automotive, and OEM/Other. Data Center has grown from a minority of revenue five years ago to representing ~90% of the FY2026 total, driven by demand for H100, H200, and Blackwell GPUs used in AI model training and inference workloads.

### Revenue Segmentation (FY2026A — $215.9B Total)

| SEGMENT | REVENUE ($B) | % OF TOTAL | YOY GROWTH | KEY DRIVERS |
|---|---|---|---|---|
| **Data Center** | ~$195.0 | ~90% | +69% | H100/H200/ Blackwell GPU, HPC, Sovereign AI |
| **Gaming** | ~$13.5 | ~6% | +19% | RTX 50 series, GeForce installed base refresh |
| **Pro Visualisation** | ~$2.1 | ~1% | +11% | NVIDIA Omniverse, AEC, media |
| **Automotive** | ~$2.0 | ~1% | +18% | DRIVE Orin, robotics, autonomous systems |
| **OEM / Other** | ~$3.3 | ~2% | +12% | Edge inference, embedded modules |

### The CUDA Ecosystem

CUDA (Compute Unified Device Architecture), launched in 2006, is the primary reason NVIDIA's competitive position is structural rather than cyclical. CUDA is a parallel computing platform and API that allows developers to write GPU-accelerated software in C, C++, Python and a range of other languages. Over two decades, NVIDIA has built an ecosystem of libraries (cuDNN, cuBLAS, NCCL, TensorRT, Triton) that are deeply integrated into every major AI framework — PyTorch, JAX, TensorFlow.

Switching away from NVIDIA GPUs is not simply a hardware decision; it requires re-engineering model training and inference pipelines, retraining engineering teams, and accepting performance trade-offs. Google's TPU, whilst competitive on cost-per-token for its own workloads, requires a fundamentally different software paradigm (XLA/JAX) that is not compatible with the CUDA ecosystem. This creates high switching costs for the estimated 4 million+ developers actively using CUDA.

## 2. INVESTMENT THESIS

**PILLAR 1**
**CUDA Moat Is Structural**
Switching costs embedded in every AI framework. A decade of co-optimisation creates defensibility that hardware competitors cannot replicate in 1–2 product cycles.

**PILLAR 2**
**Blackwell / Rubin Product Cycle**
3.6M Blackwell units on backlog; sold out through mid-2026. Rubin (next-gen) production ramp underway. $500B combined backlog gives multi-year earnings visibility.

**PILLAR 3**
**Structural Hyperscaler Capex**
Google ($180B), Amazon ($200B), Microsoft (+80%) confirmed 2026 capex. AI treated as critical national and corporate infrastructure — not a discretionary spend.

### Pillar 1: The CUDA Moat

NVIDIA's competitive advantage is the product of a deliberate, decade-long strategy of software investment that most hardware companies failed to match. CUDA's value is not the language itself — it is the ecosystem of optimised libraries, toolchains, and developer familiarity that has compounded since 2006. Key dimensions:

- cuDNN and TensorRT: NVIDIA's deep learning library and inference optimisation framework are tuned per architecture, providing 2–4x throughput advantages on NVIDIA hardware over generic implementations.
- NIM (NVIDIA Inference Microservices): Launched 2024, these pre-packaged, CUDA-optimised AI models allow rapid deployment across any NVIDIA-powered cloud or on-premise environment.
- NCCL (NVIDIA Collective Communications Library): Used in every major distributed training framework. Switching to alternative hardware requires replacing NCCL with an unproven equivalent, risking significant training efficiency losses.

- Developer flywheel: An estimated 4 million+ active CUDA developers globally means that the default starting assumption for any new AI project is NVIDIA. This creates a demand-side lock-in that is qualitatively different from standard vendor relationships.

The most credible threat to the CUDA moat is not hardware parity — it is the standardisation of AI compute at a higher abstraction layer (e.g., OpenAI Triton, MLIR). We view this as a 5–10 year transition risk, not a near-term catalyst.

### Pillar 2: Blackwell Sold Out; Rubin in Production

NVIDIA's Blackwell architecture (B100, B200, GB200 NVL72 rack-scale systems) represents a step-change in training throughput and inference efficiency. Key data points that underpin our demand view:

- 3.6 million Blackwell GPUs on confirmed backlog from hyperscalers, sold out through at least mid-2026.
- Blackwell GB300 (the refreshed variant) now represents ~66% of Blackwell revenue, indicating product maturity and accelerated ramp.
- Meta has committed to deploying 1.3 million Blackwell units in 2025–2026 alone, making it NVIDIA's single largest Data Center customer.
- $500B combined Blackwell-Rubin revenue guided by CEO Jensen Huang for FY2026–FY2027 combined — a figure that would represent the largest product cycle ramp in semiconductor history.

Rubin, NVIDIA's next architecture, is in production ramp. Rubin introduces HBM4 memory and a new interconnect fabric, with performance-per-watt improvements estimated at 30–40% over Blackwell. The transition to Rubin positions FY2027E as another step-up year even as Blackwell revenue moderates.

### Pillar 3: Hyperscaler Capex Cycle Is Structural, Not Cyclical

A central bear argument is that hyperscaler capex is unsustainable and will revert. We disagree. The capex cycle has three distinct demand layers, each reinforcing the others:

- Training compute: Large language model training continues to scale with parameters and data. GPT-5 class models require orders of magnitude more compute than GPT-3. This demand is structural.
- Inference compute: Industry analysts project inference GPU demand will exceed training demand by 4.5x by 2026. As AI is embedded in consumer and enterprise products, inference capacity scales with user adoption — a compounding, not one-time, demand driver.
- Sovereign AI: Governments from UAE to Japan to France are investing in national AI compute infrastructure. NVIDIA has secured sovereign AI contracts worth billions; this segment is largely uncorrelated with hyperscaler budget cycles.

Confirmed 2026 hyperscaler capex commitments:

| COMPANY | 2026 CAPEX GUIDANCE | YOY CHANGE | PRIMARY DRIVER |
|---|---|---|---|
| **Alphabet (Google)** | $180B | +98% vs $91B (2025) | AI compute, data centre expansion |
| **Amazon (AWS)** | $200B | +56% vs $128B (2025) | AWS AI infrastructure, Trainium + NVDA GPUs |
| **Microsoft** | ~$80B+ (AI-specific) | +80%+ YoY | Azure AI capacity, OpenAI infrastructure |
| **Meta** | >$100B total capex | +60%+ YoY | AI training clusters, 1.3M Blackwell units |

## 3. VALUATION

### Methodology Overview

We use two primary valuation approaches: (1) Discounted Cash Flow (DCF) analysis across three scenarios, and (2) Comparable Company (Comps) analysis using forward P/E and EV/EBITDA multiples. We assign lower weight to DCF given the inherent sensitivity of terminal value assumptions for a hypergrowth company; however, we are transparent that the DCF at any reasonable WACC points to meaningful overvaluation. Our BUY rating is primarily a forward-multiple call, anchored in the Rubin product cycle and FY2027 earnings inflection.

### Discounted Cash Flow Analysis

| DCF SCENARIO | WACC | TERMINAL GROWTH | FY2031E FCF ($B) | EQUITY VALUE ($B) | IMPLIED PRICE |
|---|---|---|---|---|---|
| **Bear Case (Lower growth + high WACC)** | 15.0% | 3.0% | $165.6B | $1,287B | **$52** |
| **Base Case (Base growth + 13% WACC)** | 13.0% | 4.0% | $234.8B | $2,171B | **$89** |
| **Bull Case (Higher growth + 10% WACC)** | 10.0% | 4.0% | $265.2B | $3,707B | **$151** |

Key DCF assumptions: Risk-free rate 4.50% (10Y UST), equity risk premium 5.50% (Damodaran 2026), beta 1.55x. FCF margin 50–52% based on FY2025 actuals of 46.6%, with expansion as Blackwell/Rubin margins scale. Net cash estimated at $40B.

Note: The DCF base case at $89 implies the market is pricing in either a ~7.5% WACC (consistent with treating NVIDIA as a near-monopoly infrastructure asset) or significantly higher revenue growth than our base case. Both are defensible but require explicit acceptance. We flag this as a risk.

### Comparable Company Analysis — Forward Multiple Valuation

| COMPANY | TICKER | MKT CAP ($B) | REV GROWTH NTM | NTM P/E | EV/EBITDA | FCF YIELD |
|---|---|---|---|---|---|---|
| **NVIDIA** | **NVDA** | **$4,456** | **30%** | **21.6x** | **32.0x** | **2.4%** |
| AMD | AMD | $193 | 18% | 22.1x | 22.4x | 4.1% |
| Broadcom | AVGO | $852 | 16% | 24.8x | 26.1x | 3.2% |
| TSMC | TSM | $942 | 20% | 18.9x | 14.2x | 2.8% |
| Qualcomm | QCOM | $183 | 10% | 13.6x | 10.8x | 5.9% |
| Marvell | MRVL | $88 | 25% | 28.3x | 24.5x | 1.8% |
| ASML | ASML | $289 | 12% | 27.6x | 22.3x | 1.4% |
| **Sector Median (ex-NVDA)** | — | — | **17%** | **23.5x** | **22.4x** | **3.0%** |

### Price Target Derivation

We apply a blended multiple methodology, weighting forward P/E most heavily (45%) given the company's growth profile, with EV/EBITDA as a cross-check (35%) and DCF for intrinsic value context (20%). We deliberately discount the Street consensus of $272 to reflect: (1) China revenue risk not fully priced in; (2) uncertainty around Rubin ramp timing; and (3) the fundamental reality that DCF analysis does not support current valuation at standard discount rates.

| VALUATION METHOD | MULTIPLE / ASSUMPTION | BASE METRIC (FY2027E) | IMPLIED PRICE | WEIGHT |
|---|---|---|---|---|
| Forward P/E (sector prem.) | 25.0x | Non-GAAP EPS $8.52 | $213 | 45% |

| | | | | |
|---|---|---|---|---|
| EV/EBITDA | 25.0x | EBITDA $191.5B | $196 | 35% |
| DCF (Bull, 10% WACC) | 10% WACC | FCF $140B–$235B (FY27–31) | $151 | 20% |
| **Weighted Average Price Target** | **—** | **—** | **$201** | **100%** |

12-Month Price Target: $213. We round to $213 (25x our FY2027E non-GAAP EPS of $8.52), as the forward P/E remains the most reliable anchor for near-term price action in a hypergrowth semiconductor stock. The weighted blended value of $201 provides a conservative floor.

### Bull / Base / Bear Scenarios

| SCENARIO | PRICE TARGET | UPSIDE / (DOWNSIDE) | KEY ASSUMPTION |
|---|---|---|---|
| **Bull Case** | **$270** | **+51.7%** | Rubin blowout, no China export escalation; inference TAM drives FY2027E EPS to $10.00+ |
| **Base Case** | **$213** | **+19.6%** | Blackwell demand holds, Rubin on schedule, H20 China revenue partially preserved via licensing |
| **Bear Case** | **$110** | **–38.2%** | Full China embargo + DeepSeek demand destruction; consensus EPS cuts to $6.00; multiple contracts to 18x |

## 4. KEY RISKS

We rate each risk by probability (H/M/L) and potential impact on our price target. Severity is defined as the product of both dimensions.

| RISK FACTOR | PROBABILITY | IMPACT | SEVERITY | MITIGATING FACTOR |
|---|---|---|---|---|
| **China Export Controls (H20 ban)** | High | High | **Critical** | ~13% of FY2025 revenue at risk. Trump admin selectively issuing licences; NVIDIA in discussions for ~15% revenue-sharing arrangement. Non-China demand absorbs capacity in the near term. |
| **DeepSeek / Model Efficiency** | Medium | Medium | **Elevated** | DeepSeek-V3 demonstrated competitive training with 2,048 H800s at ~$5.6M cost. Risks scaling-law assumptions. However, inference TAM expansion (~4.5x vs training) is a direct offset; more efficient models → more use cases → more inference demand. |
| **Hyperscaler Custom ASICs** | Medium | Medium-High | **Elevated** | Google TPU v7 approaches cost parity on Google-specific workloads. AWS Trainium shows 30–40% price-performance advantage on training. CUDA moat remains the principal defence; ASICs optimise for specific tasks, not general AI compute. |
| **Demand Cliff (H2 2026+)** | Low-Medium | High | **Elevated** | If hyperscaler AI monetisation disappoints, capex could revert sharply. Near-term probability is low given committed capex; the risk crystallises in 2027–2028 if inference monetisation lags. |
| **TSMC / Taiwan Geopolitical Risk** | Low | Very High | **Moderate** | ~90% of NVIDIA production relies on TSMC. A Taiwan Strait conflict would be catastrophic. Diversification to Samsung and Intel Foundry in progress but <5% of current volume. |
| **Valuation Compression** | Medium | Medium | **Moderate** | If consensus EPS estimates are revised down or macro risk appetite deteriorates, the forward multiple (21.6x) could compress to 17–18x. Our bear case assumes this scenario fully. |

## 5. FINANCIAL SUMMARY & FORECASTS

All figures in USD billions unless noted. Fiscal year ends January. FY2026A and prior are actuals; FY2027E and FY2028E are analyst estimates based on consensus inputs and our own growth assumptions.

| METRIC | FY2024A | FY2025A | FY2026A | FY2027E | FY2028E |
|---|---|---|---|---|---|
| **Revenue ($B)** | **60.9** | **130.5** | **215.9** | **280.0** | **336.0** |
| YoY Revenue Growth | +126% | +114% | +65% | +30% | +20% |
| Gross Profit ($B) | 44.3 | 98.0 | 159.8 | 204.4 | 245.3 |
| Gross Margin (%) | 72.7% | 75.0% | 74.0% | 73.0% | 73.0% |
| **EBITDA ($B)** | **32.3** | **88.2** | **147.0** | **191.5** | **228.7** |
| EBITDA Margin (%) | 53.0% | 67.6% | 68.1% | 68.4% | 68.1% |
| GAAP Operating Income ($B) | 33.0 | 80.4 | 129.5 | 175.0 | 210.0 |
| GAAP Net Income ($B) | 29.8 | 72.9 | 112.0 | 145.0 | 173.5 |
| **Non-GAAP EPS (Diluted)** | **$1.29** | **$2.94** | **$4.55** | **$8.52** | **$9.80** |
| R&D Expense ($B) | 8.7 | 12.9 | 16.2 | 20.0 | 22.8 |
| R&D as % of Revenue | 14.3% | 9.9% | 7.5% | 7.1% | 6.8% |
| **Free Cash Flow ($B)** | **27.0** | **60.9** | **108.0** | **140.0** | **171.4** |
| FCF Margin (%) | 44.3% | 46.6% | 50.0% | 50.0% | 51.0% |
| CapEx ($B) | 1.1 | 2.8 | 4.3 | 5.6 | 6.7 |
| CapEx as % of Revenue | 1.8% | 2.2% | 2.0% | 2.0% | 2.0% |

## 6. INVESTMENT CONCLUSION

**RATING: BUY | PRICE TARGET: $213 | CURRENT: $178.03 | UPSIDE: +19.6%**

NVIDIA is the closest thing to a monopoly in AI compute infrastructure that currently exists. The CUDA ecosystem, the Blackwell backlog, and the structural commitment of hyperscalers to AI investment create a business with rare visibility and pricing power. At $178, the stock is trading at 21.6x next-year earnings — arguably the cheapest it has been relative to growth prospects since the AI capex cycle began.

We are not blind to the risks. Our DCF at a 13% WACC produces $89 per share. The gap between $89 and $178 is real, and it represents the premium the market assigns to NVIDIA's monopolistic market position, software moat, and execution track record. We think that premium is partially — but not fully — justified at current levels.

Our BUY is premised on the FY2027 earnings inflection driven by the Rubin architecture ramp, continued Data Center dominance, and inference TAM expansion. If these catalysts materialise on schedule and China revenue is partially preserved, $213 is achievable within 12 months. If China is fully embargoed and Rubin is delayed, $110–130 is the realistic downside.

The honest version of our recommendation: Buy NVIDIA at current levels as a moderate overweight in a tech-tilted portfolio, with clear stop-loss logic if Q1 FY2027 (April 2026) earnings show any China revenue deterioration or Rubin ramp delays. Know what you own, size it accordingly, and don't mistake Street consensus for independent analysis.

**Disclaimer:** This analysis is prepared for educational and portfolio development purposes only. It does not constitute investment advice. All views expressed are personal to the author and do not represent any institution. Past performance and analyst price targets are not a guarantee of future results. Investors should conduct their own due diligence before making any investment decision.